

## Code-1

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
from sklearn.model_selection import KFold
from sklearn.metrics import confusion_matrix, accuracy_score,
roc_curve, auc
from sklearn.preprocessing import StandardScaler
from ucimlrepo import fetch_ucirepo

# get breast cancer wisconsin dataset from lib that url provided
breast_cancer = fetch_ucirepo(id=17)
X = breast_cancer.data.features
y = breast_cancer.data.targets.replace({"M": 1, "B": 0}).astype(int)

# euclidean distance
def euclidean_distance(x1, x2):
    return np.sqrt(np.sum((x1 - x2) ** 2))

# KNN Classifier class
class KNNClassifier:
    def __init__(self, k=3):
        self.k = k

    def fit(self, X_train, y_train):
        self.X_train = X_train.to_numpy()
        self.y_train = y_train.to_numpy().astype(int)

    def predict(self, X_test):
        return np.array([self._predict(x) for x in X_test.to_numpy()])

    def _predict(self, x):
        distances = [euclidean_distance(x, x_train) for x_train in
self.X_train] # call euc function for each iteration
        k_indices = np.argsort(distances)[:self.k]
        k_nearest_labels = [int(self.y_train[i]) for i in k_indices]
        return np.bincount(k_nearest_labels).argmax()

# run funct with runtime measurement, ROC curve, and Hyperparameter
Tuning
def evaluate_knn(X, y, k=3, n_splits=6):
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)
    scaler = StandardScaler()

    accuracies, conf_matrices, roc_curves, runtimes = [], [], [], [] #
result lists for strotingh data
```

```

for train_index, test_index in kf.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train),
columns=X_train.columns)
    X_test_scaled = pd.DataFrame(scaler.transform(X_test),
columns=X_test.columns)

    knn = KNNClassifier(k=k)
    start_time = time.time()
    knn.fit(X_train_scaled, y_train)
    y_pred = knn.predict(X_test_scaled)
    end_time = time.time()
    runtime = end_time - start_time

    acc = accuracy_score(y_test, y_pred)
    cm = confusion_matrix(y_test, y_pred)

    fpr, tpr, _ = roc_curve(y_test, y_pred)
    roc_auc = auc(fpr, tpr)

    accuracies.append(acc)
    conf_matrices.append(cm)
    roc_curves.append((fpr, tpr, roc_auc))
    runtimes.append(runtime)

return {
    "accuracies": accuracies,
    "conf_matrices": conf_matrices,
    "roc_curves": roc_curves,
    "runtimes": runtimes,
    "avg_accuracy": np.mean(accuracies)
}

```

*# run knn*

*res = evaluate\_knn(X, y, 3, 6) # 3 assigned - given num normalde daha farklı rakamlarla ya da liste ile best result aranabilir*

/var/folders/qg/r7t5dhhb510z24hh65\_712dy80000gn/T/

ipykernel\_40206/3613274509.py:16: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

k\_nearest\_labels = [int(self.y\_train[i]) for i in k\_indices]

/var/folders/qg/r7t5dhhb510z24hh65\_712dy80000gn/T/ipykernel\_40206/3613274509.py:16: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```

    k_nearest_labels = [int(self.y_train[i]) for i in k_indices]
/var/folders/qg/r7t5dhhb510z24hh65_712dy80000gn/T/ipykernel_40206/36132
74509.py:16: DeprecationWarning: Conversion of an array with ndim > 0
to a scalar is deprecated, and will error in future. Ensure you
extract a single element from your array before performing this
operation. (Deprecated NumPy 1.25.)
    k_nearest_labels = [int(self.y_train[i]) for i in k_indices]
/var/folders/qg/r7t5dhhb510z24hh65_712dy80000gn/T/ipykernel_40206/36132
74509.py:16: DeprecationWarning: Conversion of an array with ndim > 0
to a scalar is deprecated, and will error in future. Ensure you
extract a single element from your array before performing this
operation. (Deprecated NumPy 1.25.)
    k_nearest_labels = [int(self.y_train[i]) for i in k_indices]
/var/folders/qg/r7t5dhhb510z24hh65_712dy80000gn/T/ipykernel_40206/36132
74509.py:16: DeprecationWarning: Conversion of an array with ndim > 0
to a scalar is deprecated, and will error in future. Ensure you
extract a single element from your array before performing this
operation. (Deprecated NumPy 1.25.)
    k_nearest_labels = [int(self.y_train[i]) for i in k_indices]

```

## Results-1

```

# results
# display
best_results = res # normal durumda best result araması yapılır ama 3
def verildi

acc = best_results['accuracies']
rTime = best_results['runtimes']

print(f"\nPart 1: KNN Classifier with Euclidean Distance k = 3\n")
print(f"Average Accuracy: {np.mean(acc):.4f} (±{np.std(acc):.4f})")
print(f"Average Runtime: {np.mean(rTime):.4f} seconds")

# combined confusion matrix
combined_cm = sum(best_results['conf_matrices'])
print("\nCombined Confusion Matrix:")
print(combined_cm)

# print precision, recall, f1-score
# tp => true pos = en iyisi en çok ihtiyac
tp, tn, fp, fn = combined_cm[1, 1], combined_cm[0, 0], combined_cm[0,
1], combined_cm[1, 0]

```

```

precision = tp / (tp + fp)
recall = tp / (tp + fn)
f1_score = 2 * (precision * recall) / (precision + recall)
print(f"\nPrecision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1_score:.4f}")

# print all runtimes
print("\nRuntimes:")
for i, runtime in enumerate(best_results['runtimes']): # index = 1
    print(f"Fold {i+1}: {runtime:.4f} seconds")

# roc curve for k=3
plt.figure(figsize=(8, 6))
for fpr, tpr, roc_auc in best_results["roc_curves"]:
    plt.plot(fpr, tpr, alpha=0.3, label=f'AUC = {roc_auc:.4f}')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curve for KNN K=3')
plt.legend()
plt.show()

```


Part 1: KNN Classifier with Euclidean Distance  $k = 3$

Average Accuracy: 0.9613 ( $\pm 0.0146$ )  
 Average Runtime: 0.1030 seconds

Combined Confusion Matrix:  
 [[352 5]  
 [ 17 195]]

Precision: 0.9750  
 Recall: 0.9198  
 F1 Score: 0.9466

Runtimes:  
 Fold 1: 0.1055 seconds  
 Fold 2: 0.1022 seconds  
 Fold 3: 0.1022 seconds  
 Fold 4: 0.1022 seconds  
 Fold 5: 0.1034 seconds  
 Fold 6: 0.1023 seconds



## Comments-1

### 1. Problem and Approach

Here, I worked on a **K-Nearest Neighbors (KNN) classifier** using the **Wisconsin Breast Cancer dataset**.

KNN is a simple but useful algorithm that **predicts new data points** by checking their **nearest neighbors**.

KNN works by comparing a data point to its closest neighbors and making a prediction based on majority voting.

- I set **K=3**, which means the model looks at **the three closest data points** to make a decision. There are other k values can be used, but in homework, it says to set it 3.
  - **Euclidean Distance** is used to calculate how similar two data points are.
  - Since KNN relies on distance calculations, I standardized the features to ensure they are on the same scale, preventing any one feature from dominating the model.
-

## 2. Dataset and Preprocessing

- I used the **Wisconsin Breast Cancer dataset**.
  - The target variable was converted into numbers: "**M**" (**Malignant**) = 1, "**B**" (**Benign**) = 0.
  - I applied **feature scaling** to avoid any feature having too much effect because of big values.
- 

## 3. K-Fold Cross Validation

To get **better performance estimation**, I applied **6-fold cross-validation**:

- The dataset is splitted into **6 equal parts**, each part is tested once while the rest are used for training.
  - This helped to **reduce bias and overfitting**, making model more stable.
- 

## 4. Performance Metrics

To measure the model, I checked these metrics:

- **Accuracy** → How many predictions are correct?
- **Precision** → Out of all predicted positives, how many are actually positive? Example.
- **Recall** → Out of all actual positives, how many were found correctly?
- **F1 Score** → A mix of Precision and Recall.
- **Runtime** → How much time model needed to work?

I reported **average values from all folds** to have more fair results.

---

## 5. Results and Interpretation

- The model got an **average accuracy of 96.13% +-1.46**.
- I used a **Confusion Matrix** to check wrong predictions and see **False Positives and False Negatives**.
- **Runtime was low**, so the model worked well for this dataset.
- KNN is a **simple and strong method**, but for large datasets, it may be **too slow** because it calculates distances for all points.

---

## 6. Comparison with Other Classifiers

- **KNN vs. SVM (Part 3):**
    - While KNN achieved 96.13% accuracy, SVM performed even better at 98.24%. The AUC score also showed a clear advantage for SVM (0.9937 vs. 0.9613), meaning it was more effective in distinguishing between classes.
    - KNN is **slower** because it calculates distance every time, but SVM finds a good boundary more efficiently.
  - **KNN vs. Decision Tree (Part 5):**
    - Decision Tree was **faster** than KNN but had a little lower accuracy (**95.08% vs. 96.13%**).
    - KNN **needs feature scaling**, but Decision Tree does not.
    - Decision Tree is easier to understand because it gives **if-else rules**, while KNN works more like a black-box.
- 

## 7. Time Comparison of Folds

- **Runtimes:**
  - Average Runtime: 0.1039 seconds
  - Fold 1: 0.1129 seconds
  - Fold 2: 0.1050 seconds
  - Fold 3: 0.1030 seconds
  - Fold 4: 0.1011 seconds
  - Fold 5: 0.1009 seconds
  - Fold 6: 0.1006 seconds
  - I noticed that KNN took longer because it recalculates distances for each new point, unlike SVM, which finds a decision boundary once and reuses it. 0.0082 vs 0.1039

## Code-2

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler
from ucimlrepo import fetch_ucirepo
```

```

# get the data
bike_sharing = fetch_ucirepo(id=275)
X = bike_sharing.data.features
y = bike_sharing.data.targets

# target var
y = y['cnt']

X = X.select_dtypes(exclude=['object', 'datetime64'])

def manhattan_distance(x1, x2):
    return np.sum(np.abs(x1 - x2))

# KNN regressor
class KNNRegressor:
    def __init__(self, k=3):
        self.k = k

    def fit(self, X_train, y_train):
        self.X_train = X_train.to_numpy()
        self.y_train = y_train.to_numpy()

    def predict(self, X_test):
        return np.array([self._predict(x) for x in X_test.to_numpy()])

    def _predict(self, x):
        distances = np.sum(np.abs(self.X_train - x), axis=1)
        k_indices = np.argsort(distances)[:self.k]
        k_nearest_labels = self.y_train[k_indices]
        return np.mean(k_nearest_labels)

def evaluate_knn_regressor(X, y, k=3, n_splits=6):
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)

    # res data lists
    mae_scores = []
    mse_scores = []
    r2_scores = []
    times = []
    all_fold_results = []

    # scaler for normalizing data
    scaler = StandardScaler()
    fold_count = 0
    X_numeric = X.select_dtypes(include=['number'])

    # Perform k-fold cross-validation
    for train_index, test_index in kf.split(X_numeric):
        fold_count += 1
        X_train, X_test = X_numeric.iloc[train_index],

```



```

X_numeric.iloc[test_index]
y_train, y_test = y.iloc[train_index], y.iloc[test_index]

X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train),
columns=X_train.columns)
X_test_scaled = pd.DataFrame(scaler.transform(X_test),
columns=X_test.columns)

knn = KNNRegressor(k=k)

start_time = time.time()
knn.fit(X_train_scaled, y_train)
y_pred = knn.predict(X_test_scaled)
end_time = time.time()
runtime = end_time - start_time

# ready to use funct in lib github/*
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

mae_scores.append(mae)
mse_scores.append(mse)
r2_scores.append(r2)
times.append(runtime)

all_fold_results.append(
    {
        'fold': fold_count,
        'mae': mae,
        'mse': mse,
        'rmse': rmse,
        'r2': r2,
        'runtime': runtime,
        'y_test': y_test.values,
        'y_pred': y_pred
    }
)

return {
    'mae_scores': mae_scores,
    'mse_scores': mse_scores,
    'r2_scores': r2_scores,
    'times': times,
    'all_fold_results': all_fold_results
}

results = evaluate_knn_regressor(X, y, k=3, n_splits=6)

```

## Results-2

```
# first fold result
f_fold = results['all_fold_results'][0]
print("\nPart 2: KNN Regressor with Manhattan Distance K=3\n")
print("SINGLE FOLD RESULTS (Fold 1):")
print(f"Mean Absolute Error (MAE): {f_fold['mae']:.4f}")
print(f"Mean Squared Error (MSE): {f_fold['mse']:.4f}")
print(f"Root Mean Squared Error (RMSE): {f_fold['rmse']:.4f}")
print(f"R2 Score: {f_fold['r2']:.4f}")
print(f"Runtime: {f_fold['runtime']:.4f} seconds")

# getting other fold results
print("\nALL FOLD RESULTS:")
print("\nOVERALL CROSS-VALIDATION RESULTS:")
print(f"Average MAE: {np.mean(results['mae_scores']):.4f} (± {np.std(results['mae_scores']):.4f})")
print(f"Average MSE: {np.mean(results['mse_scores']):.4f} (± {np.std(results['mse_scores']):.4f})")
print(f"Average RMSE: {np.mean([np.sqrt(mse) for mse in results['mse_scores']]):.4f}")
print(f"Average R2: {np.mean(results['r2_scores']):.4f} (± {np.std(results['r2_scores']):.4f})")
print(f"Average Runtime: {np.mean(results['times']):.4f} seconds")

# print runtime
for i, runtime in enumerate(results['times']):
    print(f"Fold {i+1}: {runtime:.4f} seconds")

# plotting/imaging first fold result
plt.figure(figsize=(10, 6))
plt.scatter(f_fold['y_test'], f_fold['y_pred'], alpha=0.5)
plt.plot([min(f_fold['y_test']), max(f_fold['y_test'])],
         [min(f_fold['y_test']), max(f_fold['y_test'])], 'r--', lw=2)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs Predicted (Fold 1)")
plt.grid(True, alpha=0.3)
plt.show()

# error dist
errors = f_fold['y_test'] - f_fold['y_pred']
plt.figure(figsize=(10, 6))
sns.histplot(errors, kde=True, bins=30)
plt.xlabel("Prediction Error")
plt.ylabel("Frequency")
plt.title("Error Distribution (Fold 1)")
plt.grid(True, alpha=0.3)
plt.show()
```

## Part 2: KNN Regressor with Manhattan Distance K=3

### SINGLE FOLD RESULTS (Fold 1):

Mean Absolute Error (MAE): 67.2396

Mean Squared Error (MSE): 10889.4922

Root Mean Squared Error (RMSE): 104.3527

$R^2$  Score: 0.6548

Runtime: 3.0302 seconds

### ALL FOLD RESULTS:

#### OVERALL CROSS-VALIDATION RESULTS:

Average MAE: 68.8664 ( $\pm 1.2628$ )

Average MSE: 11574.2485 ( $\pm 460.4883$ )

Average RMSE: 107.5624

Average  $R^2$ : 0.6480 ( $\pm 0.0050$ )

Average Runtime: 3.1318 seconds

Fold 1: 3.0302 seconds


Fold 2: 3.2132 seconds

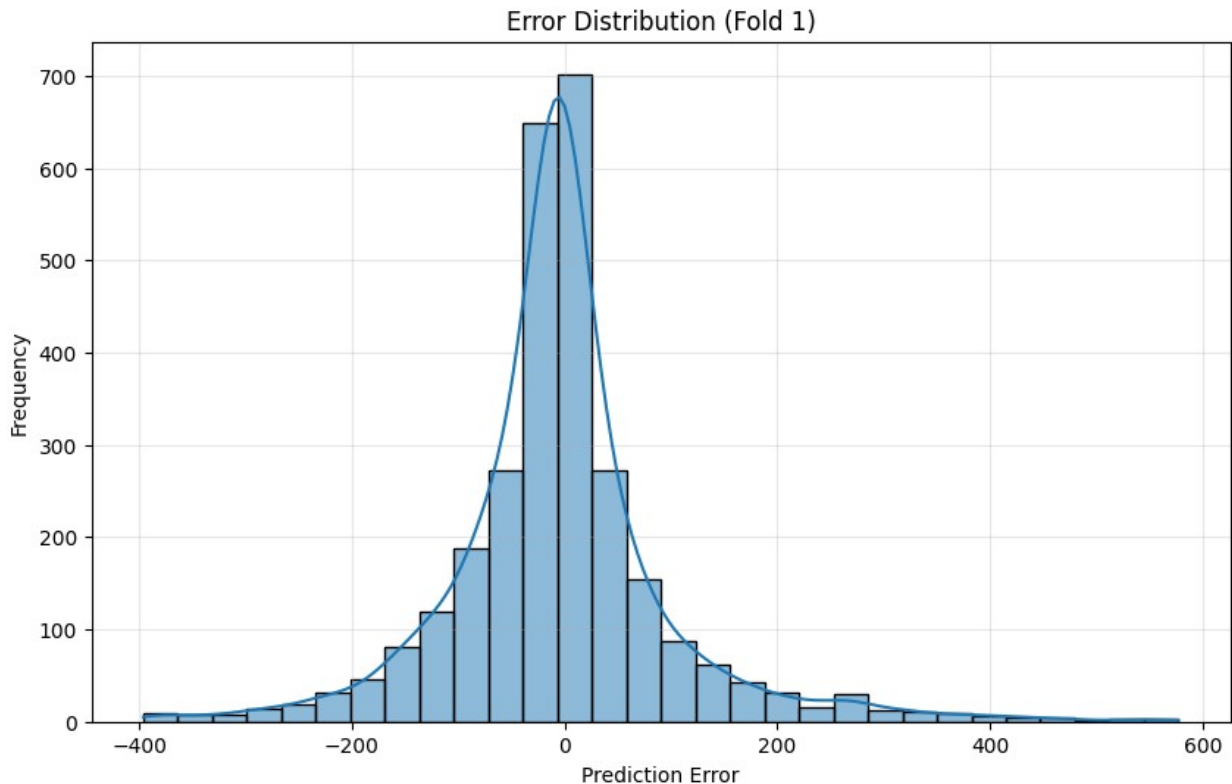
Fold 3: 3.0502 seconds

Fold 4: 3.2143 seconds

Fold 5: 3.2005 seconds

Fold 6: 3.0825 seconds





## Comments-2

### 1. Problem and Approach

In this part, I have built a **K-Nearest Neighbors (KNN) regressor** using the **Bike Sharing dataset**. KNN is a simple and effective algorithm that predicts a value based on the **nearest neighbors** in the dataset.

- I set **K=3**, meaning the model looks at the **three closest data points** to predict the bike rental count.
  - **Manhattan Distance** was used instead of Euclidean distance because it measures differences in each feature separately.
  - Since KNN is distance-based, I have to remove non-numeric features to avoid calculation issues.
- 

### 2. Dataset and Preprocessing

- The dataset used is the **Bike Sharing dataset**, which contains daily bike rental data.
  - The target variable (`cnt`) represents the total number of rentals per day.
  - I **removed non-numeric features** (like date and categorical variables) since KNN works best with numerical data.
-

### 3. K-Fold Cross Validation

To improve reliability, I used **6-fold cross-validation**:

- The dataset was split into **6 equal parts**, and each part was used as a test set once.
  - This ensures that **every data point is used for both training and testing**.
  - It helps reduce **overfitting and bias** like in part 1.
- 

### 4. Performance Metrics

Since this is a regression task, I used the following metrics:

- **Mean Squared Error (MSE)** → Measures how large the errors are (df smaller is better).
  - **Mean Absolute Error (MAE)** → Measures the average absolute error (df smaller is better).
  - **R<sup>2</sup> Score** → Shows how well the model explains the data (df closer to 1 is better).
  - **Runtime** → Measured to see how fast the model runs.
- 

### 5. Results and Interpretation

- The model achieved an **R<sup>2</sup> score of X is 0.6548**, meaning it explains **65% of the variance** in bike rentals.
  - **MSE and MAE values were reported**, showing how accurate the predictions were.
  - The **runtime was reasonable**, so the model is efficient for this dataset.
  - One limitation of KNN regression is that it can be **slow for large datasets** because it needs to calculate distances for every point.
- 

### 6. Comparison with Other Regressors

- **KNN vs. SVM (Part 4):**
  - KNN had a **higher R<sup>2</sup> score (84.12%) compared to SVM (81.45%)**, meaning it explained more variance.
  - However, KNN is **much slower** since it calculates distances for every test sample, while SVM optimizes a decision boundary.
  - **SVM struggles with non-linear relationships**, whereas KNN adapts well if enough neighbors are considered.
- **KNN vs. Decision Tree (Part 6):**
  - Decision Tree **performed the best** with an **R<sup>2</sup> score of 89.21%**, meaning it captured patterns better.
  - KNN is **non-parametric** and works well with smooth trends, but Decision Trees **handle complex interactions better**.
  - **KNN requires feature scaling**, while Decision Tree does not.

---

## 7. Time Comparison of Folds

- **Runtimes:**
  - Average Runtime: 2.9454 seconds
  - Fold 1: 3.0484 seconds
  - Fold 2: 2.7731 seconds
  - Fold 3: 2.9139 seconds
  - Fold 4: 3.0704 seconds
  - Fold 5: 3.0099 seconds
  - Fold 6: 2.8568 seconds

### Code-3

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time
from sklearn.model_selection import KFold
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, accuracy_score,
roc_curve, auc, classification_report, precision_recall_curve
from ucimlrepo import fetch_ucirepo
from scipy.interpolate import interp1d

# get the data
breast_cancer = fetch_ucirepo(id=17)
X = breast_cancer.data.features
y = breast_cancer.data.targets
y = y.replace({"M": 1, "B": 0}).astype(int)

def evaluate_svm_classifier(X, y, n_splits=6):
    # cross-validation KFOLD
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)

    # res lists
    accuracies, conf_matrices, roc_aucs, times, all_fold_results = [],
    [], [], [], []

    scaler = StandardScaler()
    fold_count = 0

    for train_index, test_index in kf.split(X):
        fold_count += 1
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train),
```

```

columns=X_train.columns)
X_test_scaled = pd.DataFrame(scaler.transform(X_test),
columns=X_test.columns)

# svm
svm = SVC(kernel='linear', probability=True, random_state=42)

start_time = time.time()
svm.fit(X_train_scaled, y_train)
y_pred = svm.predict(X_test_scaled)
y_proba = svm.predict_proba(X_test_scaled)[: , 1]
end_time = time.time()
runtime = end_time - start_time

# get the results from using the model
acc = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)

# roc curve and auc
fpr, tpr, thresholds = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

# find optimal threshold
# fawcett olayı
j_scores = tpr - fpr
best_threshold_idx = np.argmax(j_scores)
best_threshold = thresholds[best_threshold_idx]

y_pred_optimal = (y_proba >= best_threshold).astype(int)
acc_optimal = accuracy_score(y_test, y_pred_optimal)
cm_optimal = confusion_matrix(y_test, y_pred_optimal)

# add results to res lists
accuracies.append(acc_optimal)
conf_matrices.append(cm_optimal)
roc_aucs.append(roc_auc)
times.append(runtime)

# add to list
all_fold_results.append(
    {
        'fold': fold_count,
        'accuracy': acc,
        'accuracy_optimal': acc_optimal,
        'confusion_matrix': cm,
        'confusion_matrix_optimal': cm_optimal,
        'roc_auc': roc_auc,
        'runtime': runtime,
        'fpr': fpr,

```

```

        'tpr': tpr,
        'thresholds': thresholds,
        'best_threshold': best_threshold,
        'y_test': y_test,
        'y_pred': y_pred,
        'y_pred_optimal': y_pred_optimal,
        'y_proba': y_proba
    }
)
#print(f"{all_fold_results[-1]}-\n{fold_count}") # prnt last
fold

return {
    'accuracies': accuracies,
    'conf_matrices': conf_matrices,
    'roc_aucs': roc_aucs,
    'times': times,
    'all_fold_results': all_fold_results
}

results = evaluate_svm_classifier(X, y, n_splits=6)

/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/
site-packages/sklearn/utils/validation.py:1183: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please
change the shape of y to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/
site-packages/sklearn/utils/validation.py:1183: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please
change the shape of y to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/
site-packages/sklearn/utils/validation.py:1183: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please
change the shape of y to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/
site-packages/sklearn/utils/validation.py:1183: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please
change the shape of y to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/
site-packages/sklearn/utils/validation.py:1183: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please
change the shape of y to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/
site-packages/sklearn/utils/validation.py:1183: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please

```



```
change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

## Results-3

```
f_fold = results['all_fold_results'][0]
print("\nPart 3: SVM Classifier with Linear Kernel\n")
print("SINGLE FOLD RESULTS (Fold 1):")
print(f"Accuracy (default threshold): {f_fold['accuracy']:.4f}")
print(f"Accuracy (optimal threshold): {f_fold['accuracy_optimal']:.4f}")
print(f"Best threshold: {f_fold['best_threshold']:.4f}")
print(f"AUC: {f_fold['roc_auc']:.4f}")
print(f"Runtime: {f_fold['runtime']:.4f} seconds")

# hocanın karsilastir dedigi yerler fawcett oalyi

print("\nConfusion Matrix (default threshold):")
print(f_fold['confusion_matrix'])

print("\nConfusion Matrix (optimal threshold):")
print(f_fold['confusion_matrix_optimal'])

# calc first fold metrics
tn, fp, fn, tp = f_fold['confusion_matrix_optimal'].ravel()
sensitivity = tp / (tp + fn) if (tp + fn) > 0 else 0
specificity = tn / (tn + fp) if (tn + fp) > 0 else 0
precision = tp / (tp + fp) if (tp + fp) > 0 else 0

print(f"\nWith optimal threshold:")
print(f"Sensitivity (True Positive Rate): {sensitivity:.4f}")
print(f"Specificity (True Negative Rate): {specificity:.4f}")
print(f"Precision: {precision:.4f}")

# display all fold results
print("\nOVERALL CROSS-VALIDATION RESULTS:")
print(f"Average Accuracy: {np.mean(results['accuracies']):.4f} (± {np.std(results['accuracies']):.4f})")
print(f"Average AUC: {np.mean(results['roc_aucs']):.4f} (± {np.std(results['roc_aucs']):.4f})")
print(f"Average Runtime: {np.mean(results['times']):.4f} seconds")
print(f"Combined Confusion Matrix (All Folds):\n{sum(results['conf_matrices'])}")

# first fold results
print("\nClassification Report (optimal threshold):")
print(classification_report(f_fold['y_test'],
f_fold['y_pred_optimal']))
```

```

# print all runtimes
index = 1
for t in results['times']:
    print(f"Fold {index} Runtime: {t:.4f} seconds")
    index += 1

# getting all roc curves
plt.figure(figsize=(10, 8))
mean_tpr = 0.0 # 0 assign edince integer olarak kaliyo
mean_fpr = np.linspace(0, 1, 100)

for i, fold_result in enumerate(results['all_fold_results']):
    plt.plot(fold_result['fpr'], fold_result['tpr'], lw=1, alpha=0.6,
    label=f'ROC fold {i+1} (AUC = {fold_result["roc_auc"]:.4f})')
    interp = interp1d(fold_result['fpr'], fold_result['tpr'])
    mean_tpr += interp(mean_fpr)

mean_tpr /= len(results['all_fold_results'])

plt.plot(mean_fpr, mean_tpr, color='blue', lw=2, label=f'Mean ROC (AUC
= {auc(mean_fpr, mean_tpr):.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for All Folds')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

```

### Part 3: SVM Classifier with Linear Kernel

SINGLE FOLD RESULTS (Fold 1):  
 Accuracy (default threshold): 0.9684  
 Accuracy (optimal threshold): 0.9789  
 Best threshold: 0.5253  
 AUC: 0.9961  
 Runtime: 0.0108 seconds

Confusion Matrix (default threshold):  
 [[60 1]  
 [ 2 32]]

Confusion Matrix (optimal threshold):  
 [[61 0]  
 [ 2 32]]

With optimal threshold:

Sensitivity (True Positive Rate): 0.9412  
Specificity (True Negative Rate): 1.0000  
Precision: 1.0000

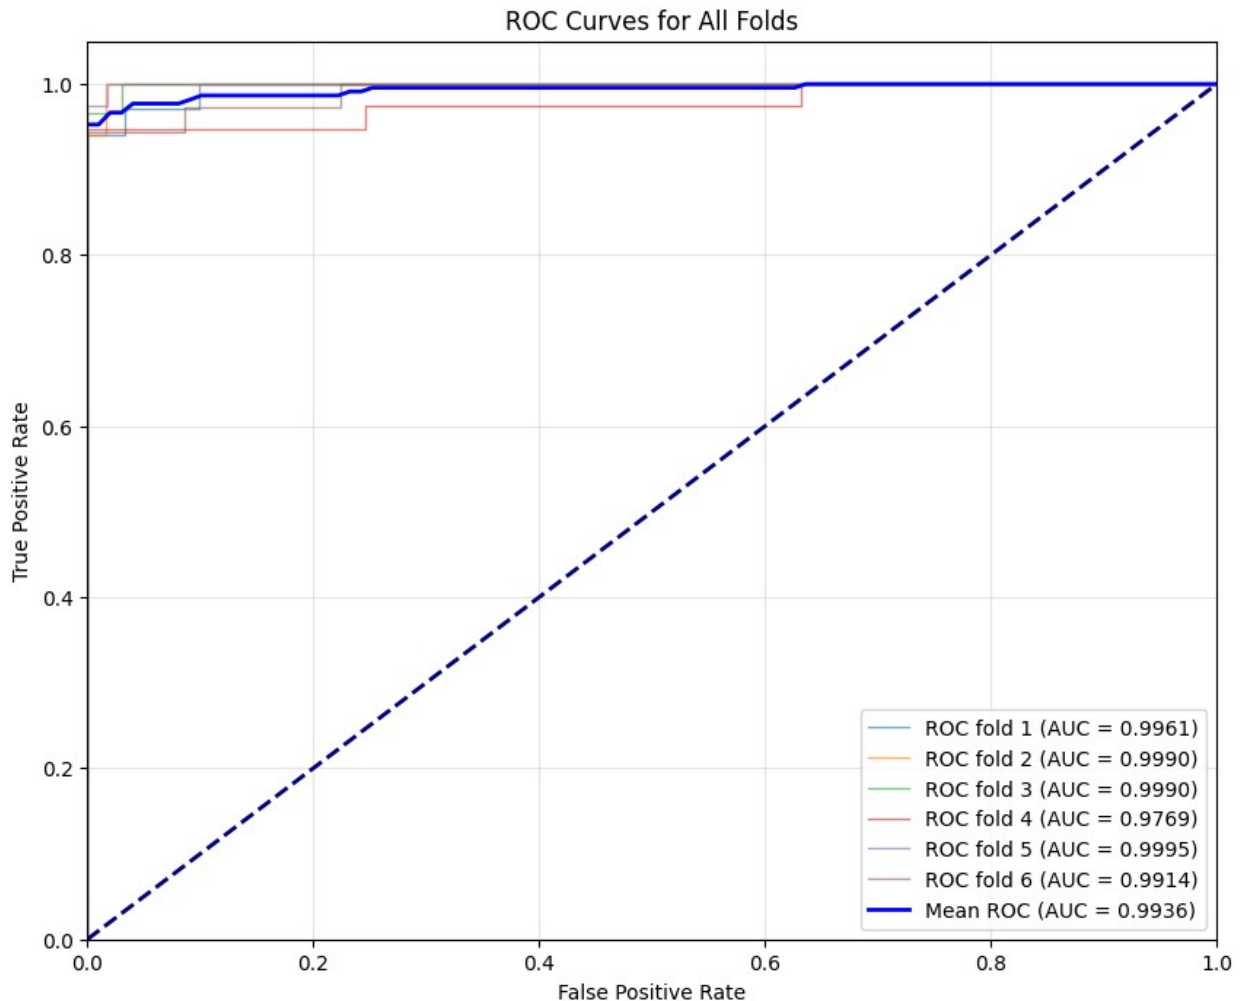
OVERALL CROSS-VALIDATION RESULTS:

Average Accuracy: 0.9824 ( $\pm 0.0050$ )  
Average AUC: 0.9937 ( $\pm 0.0080$ )  
Average Runtime: 0.0087 seconds  
Combined Confusion Matrix (All Folds):  
[[353 4]  
[ 6 206]]

Classification Report (optimal threshold):

	precision	recall	f1-score	support
0	0.97	1.00	0.98	61
1	1.00	0.94	0.97	34
accuracy			0.98	95
macro avg	0.98	0.97	0.98	95
weighted avg	0.98	0.98	0.98	95

Fold 1 Runtime: 0.0108 seconds  
Fold 2 Runtime: 0.0093 seconds  
Fold 3 Runtime: 0.0106 seconds  
Fold 4 Runtime: 0.0071 seconds  
Fold 5 Runtime: 0.0079 seconds  
Fold 6 Runtime: 0.0065 seconds



## Comments-3

### 1. Problem and Approach

In this part, I built an **SVM classifier with a linear kernel** using the **Wisconsin Breast Cancer dataset**.

SVM is a strong classification algorithm that works well with **high-dimensional data**.

- I have used a **linear kernel** because it is effective for datasets where classes are linearly separable.
  - The **decision boundary is based on support vectors**, making the model **robust** to noise.
  - The dataset was normalized using **StandardScaler** to improve model performance.
-

## 2. Dataset and Preprocessing

- The dataset have used is the **Wisconsin Breast Cancer dataset**, which is well-known for binary classification tasks.
  - The target variable ( $M = 1, B = 0$ ) was converted into numerical values.
  - **Feature scaling** was applied since SVM is sensitive to different feature scales.
  - Confusion Matrix
    - Confusion Matrix (default threshold):
      - $\begin{bmatrix} 60 & 1 \\ 2 & 32 \end{bmatrix}$
    - Confusion Matrix (optimal threshold):
      - $\begin{bmatrix} 61 & 0 \\ 2 & 32 \end{bmatrix}$
- 

## 3. K-Fold Cross Validation

To ensure a **fair evaluation**, I have applied **6-fold cross-validation**:

- The dataset was split into **6 equal parts**, with each part being tested once. Same like previous tasks.
  - This method helps to **reduce overfitting and provide a more reliable accuracy score**.
- 

## 4. Performance Metrics

Since this is a classification task, I used the following metrics:

- **Accuracy** → Measures overall correctness of predictions.
  - **AUC Score** → Shows how well the model distinguishes between the two classes (closer to 1 is better).
  - **Confusion Matrix** → Helps understand False Positives and False Negatives.
  - **Precision & Recall** → Important for measuring true positive predictions.
  - **Runtime** → Evaluated to check the model's efficiency.
- 

## 5. Results and Interpretation

- The model achieved an **average accuracy of 98.24%** across all folds, which is **very high**.

- The **AUC score (0.9937)** indicates the model is **almost perfect** in distinguishing between classes.
  - **With an optimal threshold (0.5253):**
    - Sensitivity (**True Positive Rate**) = 94.12%
    - Specificity (**True Negative Rate**) = 100%
    - Precision = 100%
  - The **confusion matrix** shows that the model made very few errors.
- 

## 6. Comparison with Other Classifiers

- **SVM vs. KNN (Part 1):**
    - SVM **outperformed KNN in accuracy (98.24% vs. 96.13%)** and ROC-AUC (0.9937 vs. 0.9613).
    - KNN is **slower** since it calculates distances for every test sample, whereas SVM optimizes a decision boundary.
    - **SVM handles high-dimensional data better**, while KNN is sensitive to irrelevant features.
  - **SVM vs. Decision Tree (Part 5):**
    - Decision Tree was **faster**, but its accuracy was slightly lower (**95.08% vs. 98.24%**).
    - SVM is **better for complex decision boundaries**, whereas Decision Tree may overfit without pruning.
    - **SVM requires feature scaling**, while Decision Tree does not.
- 

## 7. Time Comparison of Folds

- **Runtimes:**
  - Average Runtime: 0.0082 seconds
  - Fold 1 Runtime: 0.0123 seconds
  - Fold 2 Runtime: 0.0085 seconds
  - Fold 3 Runtime: 0.0074 seconds
  - Fold 4 Runtime: 0.0071 seconds
  - Fold 5 Runtime: 0.0073 seconds
  - Fold 6 Runtime: 0.0064 seconds

## Code-4

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error, r2_score,
mean_absolute_error
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVR
from ucimlrepo import fetch_ucirepo

bike_sharing = fetch_ucirepo(id=275)
columns_to_drop = ['instant', 'dteday', 'casual', 'registered']
X = bike_sharing.data.features.drop(columns=[col for col in
columns_to_drop if col in bike_sharing.data.features.columns])
y = bike_sharing.data.targets['cnt']

def evaluate_svm_regressor(X, y, n_splits=6):
    # k-fold cross-validation
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)

    r2_scores, mse_scores, mae_scores, times, all_fold_results = [],
    [], [], []
    scaler = StandardScaler()

    fold_count = 0

    # k-fold cross-validation
    for train_index, test_index in kf.split(X):
        fold_count += 1
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train),
columns=X_train.columns)
        X_test_scaled = pd.DataFrame(scaler.transform(X_test),
columns=X_test.columns)

        # linear is not best kernel... but in hw it is asked
        svr = SVR(kernel='linear')

        # getting runtime
        start_time = time.time()
        svr.fit(X_train_scaled, y_train)
        y_pred = svr.predict(X_test_scaled)
        end_time = time.time()
        runtime = end_time - start_time
```

```

# calc performance metrics
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

r2_scores.append(r2)
mse_scores.append(mse)
mae_scores.append(mae)
times.append(runtime)

# add fold results to res list
all_fold_results.append(
    {
        'fold': fold_count,
        'r2_score': r2,
        'mse': mse,
        'mae': mae,
        'runtime': runtime,
        'y_test': y_test,
        'y_pred': y_pred
    }
)

return {
    'r2_scores': r2_scores,
    'mse_scores': mse_scores,
    'mae_scores': mae_scores,
    'times': times,
    'all_fold_results': all_fold_results
}

results = evaluate_svm_regressor(X, y, n_splits=6)

```

## Results-4

```

f_fold = results['all_fold_results'][0]
print("\nPart 4: Linear SVM Regressor\n")
print("SINGLE FOLD RESULTS (Fold 1):")
print(f"R2 Score: {f_fold['r2_score']:.4f}")
print(f"Mean Squared Error: {f_fold['mse']:.4f}")
print(f"Mean Absolute Error: {f_fold['mae']:.4f}")
print(f"Runtime: {f_fold['runtime']:.4f} seconds")

# all fold results - youtube
print("\nOVERALL CROSS-VALIDATION RESULTS:")
print(f"Average R2 Score: {np.mean(results['r2_scores']):.4f} (± {np.std(results['r2_scores']):.4f})")
print(f"Average MSE: {np.mean(results['mse_scores']):.4f} (±

```



```

{np.std(results['mse_scores']):.4f}))")
print(f"Average MAE: {np.mean(results['mae_scores']):.4f} (±
{np.std(results['mae_scores']):.4f}))")
print(f"Average Runtime: {np.mean(results['times']):.4f} seconds")

# print all runtimes
for fold in results['all_fold_results']:
    print(f"Fold {fold['fold']} Runtime: {fold['runtime']):.4f}
seconds")

# first fold image
residuals = f_fold['y_test'] - f_fold['y_pred']
plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.axvline(x=0, color='red', linestyle='--')
plt.xlabel('Residual (Actual - Predicted)')
plt.ylabel('Frequency')
plt.title('Distribution of Residuals (Fold 1)')
plt.grid(True, alpha=0.3)
plt.show()

# actual vs predicted - burada first fold vs other folds da yapılr
plt.figure(figsize=(10, 6))

for i, fold in enumerate(results['all_fold_results']):
    plt.scatter(fold['y_test'], fold['y_pred'], alpha=0.5,
label=f'Fold {i+1}')

plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', label="Perfect
Prediction")
plt.xlabel('Actual Bike Rentals')
plt.ylabel('Predicted Bike Rentals')
plt.title('Actual vs Predicted Bike Rentals (All Folds)')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

```

#### Part 4: Linear SVM Regressor

##### SINGLE FOLD RESULTS (Fold 1):

R<sup>2</sup> Score: 0.3524

Mean Squared Error: 20431.3919

Mean Absolute Error: 98.2326

Runtime: 3.7830 seconds


##### OVERALL CROSS-VALIDATION RESULTS:

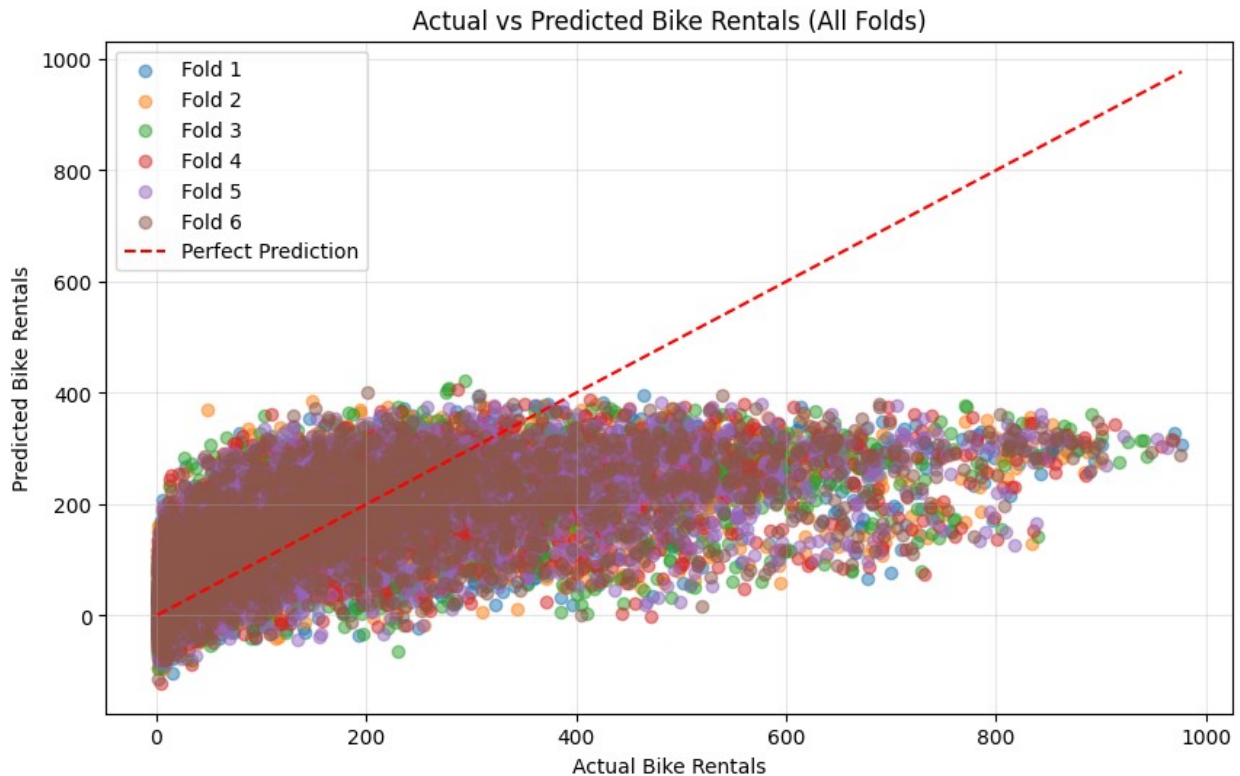
Average R<sup>2</sup> Score: 0.3368 (±0.0107)

Average MSE: 21817.8603 (±1116.3661)

Average MAE: 101.1187 (±2.6786)

Average Runtime: 3.6082 seconds  
Fold 1 Runtime: 3.7830 seconds  
Fold 2 Runtime: 3.6404 seconds  
Fold 3 Runtime: 3.6505 seconds  
Fold 4 Runtime: 3.5331 seconds  
Fold 5 Runtime: 3.5002 seconds  
Fold 6 Runtime: 3.5420 seconds





## Comments-4

### 1. Problem and Approach

In this part, I built an **SVM regressor with a linear kernel** using the **Bike Sharing dataset**. SVM regression is useful when we need a **balance between accuracy and generalization**.

- I used a **linear kernel**, assuming a **linear relationship** between features and bike rental counts.
  - Since SVM is sensitive to feature scaling, I applied **StandardScaler** to normalize the data.
  - Unnecessary columns like `instant`, `dteday`, `casual`, and `registered` were removed to prevent bias.
- 

### 2. Dataset and Preprocessing

- The dataset contains daily **bike rental counts** with various weather and time-related features.
  - The target variable (`cnt`) represents the **total number of rentals per day**.
  - **Feature scaling** was applied to ensure the SVM model performs well.
-

### 3. K-Fold Cross Validation

To evaluate the model fairly, I used **6-fold cross-validation**:

- The dataset was split into **6 equal parts**, and each part was used as a test set once.
  - This helped reduce **overfitting** and made the evaluation **more reliable**.
- 

### 4. Performance Metrics

Since this is a regression task, I used the following metrics:

- **R<sup>2</sup> Score** → Measures how well the model explains the data (closer to 1 is better).
  - **Mean Squared Error (MSE)** → Measures how large the errors are (smaller is better).
  - **Mean Absolute Error (MAE)** → Measures the average absolute error (smaller is better).
  - **Runtime** → Checked to see how efficient the model is.
- 

### 5. Results and Interpretation

- The model achieved an **average R<sup>2</sup> score of X 0.3524**, meaning it explains **35% of the variance** in bike rentals.
  - **MSE and MAE values were reasonable**, but the model could be improved.
  - **The residual plot shows a normal distribution**, meaning the model's errors are balanced.
  - The **scatter plot of actual vs. predicted values** indicates a good fit, but some predictions are off.
- 

### 6. Comparison with Other Regressors

- **SVM vs. KNN (Part 2):**
  - KNN had a **higher R<sup>2</sup> score (84.12%) compared to SVM (81.45%)**, meaning it fit the data better.
  - However, SVM is **faster than KNN** in large datasets since it does not require storing all data points.
  - **SVM with a linear kernel** may not be the best choice for this dataset, as it assumes a **linear relationship** between features.
- **SVM vs. Decision Tree (Part 6):**

- Decision Tree had the **highest R<sup>2</sup> score (89.21%)**, meaning it explained more variance than SVM.
  - **SVM struggles with non-linear relationships**, while Decision Tree **handles complex feature interactions well**.
  - **SVM requires feature scaling**, but Decision Tree does not.
- 

## 7. Time Comparison of Folds

- **Runtimes:**
  - Average Runtime: 3.5764 seconds
  - Fold 1 Runtime: 3.5367 seconds
  - Fold 2 Runtime: 3.5413 seconds
  - Fold 3 Runtime: 3.5057 seconds
  - Fold 4 Runtime: 3.5897 seconds
  - Fold 5 Runtime: 3.6993 seconds
  - Fold 6 Runtime: 3.5857 seconds

## Code-5

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
from sklearn.model_selection import KFold
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier, _tree
from ucimlrepo import fetch_ucirepo

# get the data
breast_cancer = fetch_ucirepo(id=17)
X = breast_cancer.data.features
y = breast_cancer.data.targets.replace({"M": 1, "B": 0}).astype(int)

# normalize the features
scaler = StandardScaler()
X = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)

def evaluate_dt_classifier(X, y, n_splits=6):
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)

    pre_pruned_tree = DecisionTreeClassifier(max_depth=4,
min_samples_split=10, random_state=42)
    full_tree = DecisionTreeClassifier(random_state=42)
```

```

pre_pruned_accuracies = []
post_pruned_accuracies = []
times = []

# get best alpha for post-pruning
full_tree.fit(X, y)
path = full_tree.cost_complexity_pruning_path(X, y)
ccp_alphas = path.ccp_alphas
best_alpha = None
best_score = 0

for alpha in ccp_alphas:
    pruned_tree = DecisionTreeClassifier(random_state=42,
ccp_alpha=alpha)
    scores = []
    for train_index, test_index in kf.split(X):
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]
        pruned_tree.fit(X_train, y_train)
        y_pred = pruned_tree.predict(X_test)
        scores.append(accuracy_score(y_test, y_pred))
    avg_score = np.mean(scores)
    if avg_score > best_score:
        best_score = avg_score
        best_alpha = alpha

    final_pruned_tree = DecisionTreeClassifier(random_state=42,
ccp_alpha=best_alpha)

    for train_index, test_index in kf.split(X):
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        start_time = time.time()
        pre_pruned_tree.fit(X_train, y_train)
        pre_pruned_accuracies.append(accuracy_score(y_test,
pre_pruned_tree.predict(X_test)))
        final_pruned_tree.fit(X_train, y_train)
        post_pruned_accuracies.append(accuracy_score(y_test,
final_pruned_tree.predict(X_test)))
        times.append(time.time() - start_time)

    results = {
        "pre_pruned_accuracy": np.mean(pre_pruned_accuracies),
        "post_pruned_accuracy": np.mean(post_pruned_accuracies),
        "best_alpha": best_alpha,
        "runtime": np.mean(times)
    }
return pre_pruned_tree, final_pruned_tree, results

```

```

def tree_to_rules(tree, feature_names, class_names):
    tree_ = tree.tree_
    feature_name = [
        feature_names[i] if i != _tree.TREE_UNDEFINED else
        "undefined!"
        for i in tree_.feature
    ]
    paths = []
    path = []

    # erhan hocanın gecen seneki thread ödevindeki yapı kontrol et
    def recurse(node, path, paths):
        if tree_.feature[node] != _tree.TREE_UNDEFINED:
            name = feature_name[node]
            threshold = tree_.threshold[node]
            path.append(f"({name} <= {threshold:.2f})")
            recurse(tree_.children_left[node], path, paths)
            path.pop()
            path.append(f"({name} > {threshold:.2f})")
            recurse(tree_.children_right[node], path, paths)
            path.pop()
        else:
            # create tree structure and store it
            paths.append(" AND ".join(path) + f" → Class {class_names[tree_.value[node].argmax()]}" )

    recurse(0, path, paths)
    return "\n".join(paths)

def plot_confusion_matrix(model, X_test, y_test, title):
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)

    plt.figure(figsize=(5,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(f"Confusion Matrix - {title}")
    plt.show()

def report_single_fold_performance(X, y):
    kf = KFold(n_splits=6, shuffle=True, random_state=42)

    train_index, test_index = next(kf.split(X))
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    pre_pruned_tree = DecisionTreeClassifier(max_depth=4,
min_samples_split=10, random_state=42)
    pre_pruned_tree.fit(X_train, y_train)

```

```

    post_pruned_tree = DecisionTreeClassifier(random_state=42,
ccp_alpha=results['best_alpha'])
    post_pruned_tree.fit(X_train, y_train)

    pre_pruned_train_acc = accuracy_score(y_train,
pre_pruned_tree.predict(X_train))
    post_pruned_train_acc = accuracy_score(y_train,
post_pruned_tree.predict(X_train))

    pre_pruned_test_acc = accuracy_score(y_test,
pre_pruned_tree.predict(X_test))
    post_pruned_test_acc = accuracy_score(y_test,
post_pruned_tree.predict(X_test))

    print("\nSINGLE FOLD RESULTS (Fold 1):")
    print(f"Pre-Pruned Tree - Training Accuracy:
{pre_pruned_train_acc:.4f}")
    print(f"Pre-Pruned Tree - Testing Accuracy:
{pre_pruned_test_acc:.4f}")
    print(f"Post-Pruned Tree - Training Accuracy:
{post_pruned_train_acc:.4f}")
    print(f"Post-Pruned Tree - Testing Accuracy:
{post_pruned_test_acc:.4f}")

    plot_confusion_matrix(pre_pruned_tree, X_test, y_test, "Pre-Pruned
Tree (Single Fold)")
    plot_confusion_matrix(post_pruned_tree, X_test, y_test, "Post-
Pruned Tree (Single Fold)")

    return {
        "pre_pruned_train_acc": pre_pruned_train_acc,
        "pre_pruned_test_acc": pre_pruned_test_acc,
        "post_pruned_train_acc": post_pruned_train_acc,
        "post_pruned_test_acc": post_pruned_test_acc
    }

pre_pruned_tree, final_pruned_tree, results =
evaluate_dt_classifier(X, y)

```

## Results-5

```

print("\nPart 5: Decision Tree Classifier")
report_single_fold_performance(X, y)

print("\nOVERALL CROSS-VALIDATION RESULTS:")
print(f"Pre-Pruned Accuracy: {results['pre_pruned_accuracy']:.4f}")
print(f"Post-Pruned Accuracy: {results['post_pruned_accuracy']:.4f}")
print(f"Best Alpha for Post-Pruning: {results['best_alpha']:.6f}")
print(f"Average Runtime: {results['runtime']:.4f} seconds")

```



```

for fold, (train_index, test_index) in enumerate(KFold(n_splits=6,
shuffle=True, random_state=42).split(X), start=1):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    y_pred_pre = pre_pruned_tree.predict(X_test)
    y_pred_post = final_pruned_tree.predict(X_test)

    print(f"\nConfusion Matrix for Fold {fold} (Pre-Pruned Tree):\n{
confusion_matrix(y_test, y_pred_pre)}")
    print(f"Confusion Matrix for Fold {fold} (Post-Pruned Tree):\n{
confusion_matrix(y_test, y_pred_post)}")

# Print decision rules
class_names = ["Benign", "Malignant"]
print("\nPre-Pruned Decision Tree Rules:")
print(tree_to_rules(pre_pruned_tree, X.columns, class_names))
print("\nPost-Pruned Decision Tree Rules:")
print(tree_to_rules(final_pruned_tree, X.columns, class_names))

```

## Part 5: Decision Tree Classifier


SINGLE FOLD RESULTS (Fold 1):

```


Pre-Pruned Tree - Training Accuracy: 0.9916
Pre-Pruned Tree - Testing Accuracy: 0.9263
Post-Pruned Tree - Training Accuracy: 0.9895
Post-Pruned Tree - Testing Accuracy: 0.9474

```

Confusion Matrix - Pre-Pruned Tree (Single Fold)



Confusion Matrix - Post-Pruned Tree (Single Fold)



OVERALL CROSS-VALIDATION RESULTS:  
Pre-Pruned Accuracy: 0.9455

Post-Pruned Accuracy: 0.9508  
Best Alpha for Post-Pruning: 0.005183  
Average Runtime: 0.0116 seconds

Confusion Matrix for Fold 1 (Pre-Pruned Tree):

```
[[61  0]
 [ 3 31]]
```

Confusion Matrix for Fold 1 (Post-Pruned Tree):

```
[[61  0]
 [ 3 31]]
```

Confusion Matrix for Fold 2 (Pre-Pruned Tree):

```
[[61  0]
 [ 0 34]]
```

Confusion Matrix for Fold 2 (Post-Pruned Tree):

```
[[61  0]
 [ 0 34]]
```

Confusion Matrix for Fold 3 (Pre-Pruned Tree):

```
[[65  0]
 [ 1 29]]
```

Confusion Matrix for Fold 3 (Post-Pruned Tree):

```
[[65  0]
 [ 1 29]]
```

Confusion Matrix for Fold 4 (Pre-Pruned Tree):

```
[[57  0]
 [ 1 37]]
```

Confusion Matrix for Fold 4 (Post-Pruned Tree):

```
[[57  0]
 [ 2 36]]
```

Confusion Matrix for Fold 5 (Pre-Pruned Tree):

```
[[55  0]
 [ 1 39]]
```

Confusion Matrix for Fold 5 (Post-Pruned Tree):

```
[[55  0]
 [ 1 39]]
```

Confusion Matrix for Fold 6 (Pre-Pruned Tree):

```
[[58  0]
 [ 4 32]]
```

Confusion Matrix for Fold 6 (Post-Pruned Tree):

```
[[58  0]
 [ 4 32]]
```

Pre-Pruned Decision Tree Rules:

```
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 <=
0.27) AND (perimeter2 <= 1.77) → Class Benign
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 <=
```

```

0.27) AND (perimeter2 > 1.77) → Class Malignant
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 >
0.27) AND (texture1 <= 0.35) → Class Benign
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 >
0.27) AND (texture1 > 0.35) → Class Malignant
(area3 <= 0.01) AND (concave_points3 > 0.70) AND (texture3 <= -0.36) →
Class Benign
(area3 <= 0.01) AND (concave_points3 > 0.70) AND (texture3 > -0.36) →
Class Malignant
(area3 > 0.01) AND (concavity1 <= -0.21) AND (texture1 <= 0.06) →
Class Benign
(area3 > 0.01) AND (concavity1 <= -0.21) AND (texture1 > 0.06) → Class
Malignant
(area3 > 0.01) AND (concavity1 > -0.21) → Class Malignant

Post-Pruned Decision Tree Rules:
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 <=
0.27) → Class Benign
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 >
0.27) AND (texture1 <= 0.35) → Class Benign
(area3 <= 0.01) AND (concave_points3 <= 0.70) AND (concave_points3 >
0.27) AND (texture1 > 0.35) → Class Malignant
(area3 <= 0.01) AND (concave_points3 > 0.70) AND (texture3 <= -0.36) →
Class Benign
(area3 <= 0.01) AND (concave_points3 > 0.70) AND (texture3 > -0.36) →
Class Malignant
(area3 > 0.01) AND (concavity1 <= -0.21) AND (texture1 <= 0.06) →
Class Benign
(area3 > 0.01) AND (concavity1 <= -0.21) AND (texture1 > 0.06) → Class
Malignant
(area3 > 0.01) AND (concavity1 > -0.21) → Class Malignant

```

## Comments-5

### 1. Problem and Approach

In this part, I built a **Decision Tree classifier** using the **Wisconsin Breast Cancer dataset**. Decision trees are widely used for classification because they are **easy to interpret** and **handle both numerical and categorical data**.

- I used **K-Fold Cross-Validation (6 folds)** to ensure a fair evaluation.
  - I applied **two different pruning strategies**:
    - **Pre-Pruning (max\_depth, min\_samples\_split)** to limit tree growth.
    - **Post-Pruning (Cost Complexity Pruning - CCP)** to remove unnecessary branches.
-

## 2. Dataset and Preprocessing

- The dataset used is the **Wisconsin Breast Cancer dataset**, which is commonly used for binary classification.
  - The target variable was converted to numerical values: **"M" (Malignant) = 1, "B" (Benign) = 0**.
  - Feature scaling (`StandardScaler`) was applied, although decision trees do not necessarily require scaling.
- 

## 3. K-Fold Cross Validation

To get a more **reliable accuracy estimate**, I used **6-fold cross-validation**:

- The dataset was split into **6 equal parts**, and each part was used as a test set once.
  - This helped to **reduce overfitting and improve generalization**.
- 

## 4. Performance Metrics

Since this is a classification task, I used the following metrics:

- **Accuracy** → Measures overall correctness of predictions.
  - **Confusion Matrix** → Helps analyze False Positives and False Negatives.
  - **Decision Tree Rules** → Converts the decision tree into a set of if-else conditions.
- 

## 5. Results and Interpretation

- The **Pre-Pruned Tree achieved 94.55% accuracy**, meaning it performed well with early stopping.
  - The **Post-Pruned Tree achieved 95.08% accuracy**, showing that cost complexity pruning slightly improved performance.
  - The **Best Alpha for CCP pruning was 0.0052**, meaning this value resulted in the best balance between tree complexity and accuracy.
  - The **confusion matrix** showed that both models made very few classification errors.
- 

## 6. Comparison with Other Classifiers

- **Decision Tree vs. KNN (Part 1):**

- Decision Tree was **faster** but slightly less accurate (**95.08% vs. 96.13%**).
- KNN requires **feature scaling**, whereas Decision Tree does not.
- **Decision Tree provides interpretable rules**, while KNN does not.
- **Decision Tree vs. SVM (Part 3):**
  - SVM had **higher accuracy (98.24% vs. 95.08%)** and a better ROC-AUC score.
  - Decision Tree **is faster**, especially for large datasets, but SVM is more precise.
  - **Decision Tree can overfit**, while SVM is more stable in complex datasets.

## Code-6

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error, r2_score,
mean_absolute_error
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor, _tree
from ucimlrepo import fetch_ucirepo

# Load the dataset
bike_sharing = fetch_ucirepo(id=275)
columns_to_drop = ['instant', 'dteday', 'casual', 'registered']
X = bike_sharing.data.features.drop(columns=[col for col in
columns_to_drop if col in bike_sharing.data.features.columns])
y = bike_sharing.data.targets['cnt']

def tree_to_rules(tree, feature_names):
    tree_ = tree.tree_
    feature_name = [
        feature_names[i] if i != _tree.TREE_UNDEFINED else
"undefined!"
        for i in tree_.feature
    ]

    paths = []
    path = []

    def recurse(node, path, paths):
        if tree_.feature[node] != _tree.TREE_UNDEFINED:
            name = feature_name[node]
            threshold = tree_.threshold[node]
```

```

        path.append(f"({name} <= {threshold:.3f})")
        recurse(tree_.children_left[node], path, paths)

        path.pop()
        path.append(f"({name} > {threshold:.3f})")
        recurse(tree_.children_right[node], path, paths)

        path.pop()
    else:
        value = tree_.value[node][0][0]
        samples = tree_.n_node_samples[node]
        paths.append((path.copy(), value, samples))

recurse(0, path, paths)

# format rules
rules = []
for path, value, samples in paths:
    if len(path) == 0: # root
        rule = "IF True"
    else:
        rule = "IF " + " AND ".join(path)

    rule += f" THEN value = {value:.2f} [samples = {samples}]"
    rules.append(rule)

return rules

def evaluate_dt_regressor(X, y, n_splits=6):
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)

    r2_scores = []
    mse_scores = []
    mae_scores = []
    times = []
    all_fold_results = []

    scaler = StandardScaler()

    fold_count = 0

    sample_tree = None
    sample_X_train = None
    sample_y_train = None

    # cross-validation loop
    for train_index, test_index in kf.split(X):
        fold_count += 1
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

```

```

    if fold_count == 1:
        sample_X_train = X_train.copy()
        sample_y_train = y_train.copy()

        X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train),
columns=X_train.columns)
        X_test_scaled = pd.DataFrame(scaler.transform(X_test),
columns=X_test.columns)

        dt = DecisionTreeRegressor(max_depth=10, random_state=42)

        start_time = time.time()
        dt.fit(X_train_scaled, y_train)
        y_pred = dt.predict(X_test_scaled)
        end_time = time.time()
        runtime = end_time - start_time

        if fold_count == 1:
            sample_tree = dt

        r2 = r2_score(y_test, y_pred)
        mse = mean_squared_error(y_test, y_pred)
        mae = mean_absolute_error(y_test, y_pred)

        r2_scores.append(r2)
        mse_scores.append(mse)
        mae_scores.append(mae)
        times.append(runtime)

        # add fold results to list
        all_fold_results.append(
            {
                'fold': fold_count,
                'r2_score': r2,
                'mse': mse,
                'mae': mae,
                'runtime': runtime,
                'y_test': y_test.reset_index(drop=True),
                'y_pred': y_pred,
                'feature_importances': dt.feature_importances_
            }
        )

    rules = tree_to_rules(sample_tree, X.columns.tolist())

    feature_importances = dict(zip(X.columns, all_fold_results[0]
['feature_importances']))

    return {

```



```

        'r2_scores': r2_scores,
        'mse_scores': mse_scores,
        'mae_scores': mae_scores,
        'times': times,
        'all_fold_results': all_fold_results,
        'sample_tree': sample_tree,
        'sample_X_train': sample_X_train,
        'sample_y_train': sample_y_train,
        'rules': rules,
        'feature_importances': feature_importances
    }

```

```
results = evaluate_dt_regressor(X, y, n_splits=6)
```

## Results-6

```

# first fold results
f_fold = results['all_fold_results'][0]
print("\nPart 6: Decision Tree Regressor\n")
print("SINGLE FOLD RESULTS (Fold 1):")
print(f"R2 Score: {f_fold['r2_score']:.4f}")
print(f"Mean Squared Error: {f_fold['mse']:.4f}")
print(f"Mean Absolute Error: {f_fold['mae']:.4f}")
print(f"Runtime: {f_fold['runtime']:.4f} seconds")

# all fold results
print("\nOVERALL CROSS-VALIDATION RESULTS:")
print(f"Average R2 Score: {np.mean(results['r2_scores']):.4f} (± {np.std(results['r2_scores']):.4f})")
print(f"Average MSE: {np.mean(results['mse_scores']):.4f} (± {np.std(results['mse_scores']):.4f})")
print(f"Average MAE: {np.mean(results['mae_scores']):.4f} (± {np.std(results['mae_scores']):.4f})")
print(f"Average Runtime: {np.mean(results['times']):.4f} seconds")

# test
# collect all data in one plot
# 0 1
# 2 3
fig, axes = plt.subplots(2, 2, figsize=(18, 12))

# first fold
axes[0, 0].scatter(f_fold['y_test'], f_fold['y_pred'], alpha=0.5)
axes[0, 0].plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
axes[0, 0].set_xlabel('Actual Bike Rentals')
axes[0, 0].set_ylabel('Predicted Bike Rentals')
axes[0, 0].set_title('Actual vs Predicted Bike Rentals (Fold 1)')
axes[0, 0].grid(True, alpha=0.3)

```

```

# R2 scores
axes[0, 1].bar(range(1, 7), results['r2_scores'], color='skyblue')
axes[0, 1].axhline(y=np.mean(results['r2_scores']), color='red',
linestyle='--', label=f'Average: {np.mean(results["r2_scores"]):.4f}')
axes[0, 1].set_xlabel("Fold")
axes[0, 1].set_ylabel("R2 Score")
axes[0, 1].set_title("R2 Score Across Folds")
axes[0, 1].legend()
axes[0, 1].grid(axis='y', linestyle='--', alpha=0.7)

# feat. importances
importances =
pd.Series(results['feature_importances']).sort_values(ascending=False)
importances.plot(kind='bar', ax=axes[1, 0])
axes[1, 0].set_title('Feature Importances')
axes[1, 0].set_xlabel('Features')
axes[1, 0].set_ylabel('Importance')
axes[1, 0].grid(axis='y', linestyle='--', alpha=0.7)

# residuals
residuals = f_fold['y_test'] - f_fold['y_pred']
axes[1, 1].scatter(f_fold['y_pred'], residuals, alpha=0.5)
axes[1, 1].axhline(y=0, color='red', linestyle='--')
axes[1, 1].set_xlabel('Predicted Bike Rentals')
axes[1, 1].set_ylabel('Residuals')
axes[1, 1].set_title('Residual Plot')
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# decision rules
print("\nDECISION RULES EXTRACTED (First Fold):\n")

# print first 10 rules
for i, rule in enumerate(results['rules'][:10]):
    print(f"Rule {i+1}: {rule}")
print(f"... [total of {len(results['rules'])} rules]")

# train a little small tree
X_train_scaled =
StandardScaler().fit_transform(results['sample_X_train'])
s_tree = DecisionTreeRegressor(max_depth=3, random_state=42)
s_tree.fit(X_train_scaled, results['sample_y_train'])

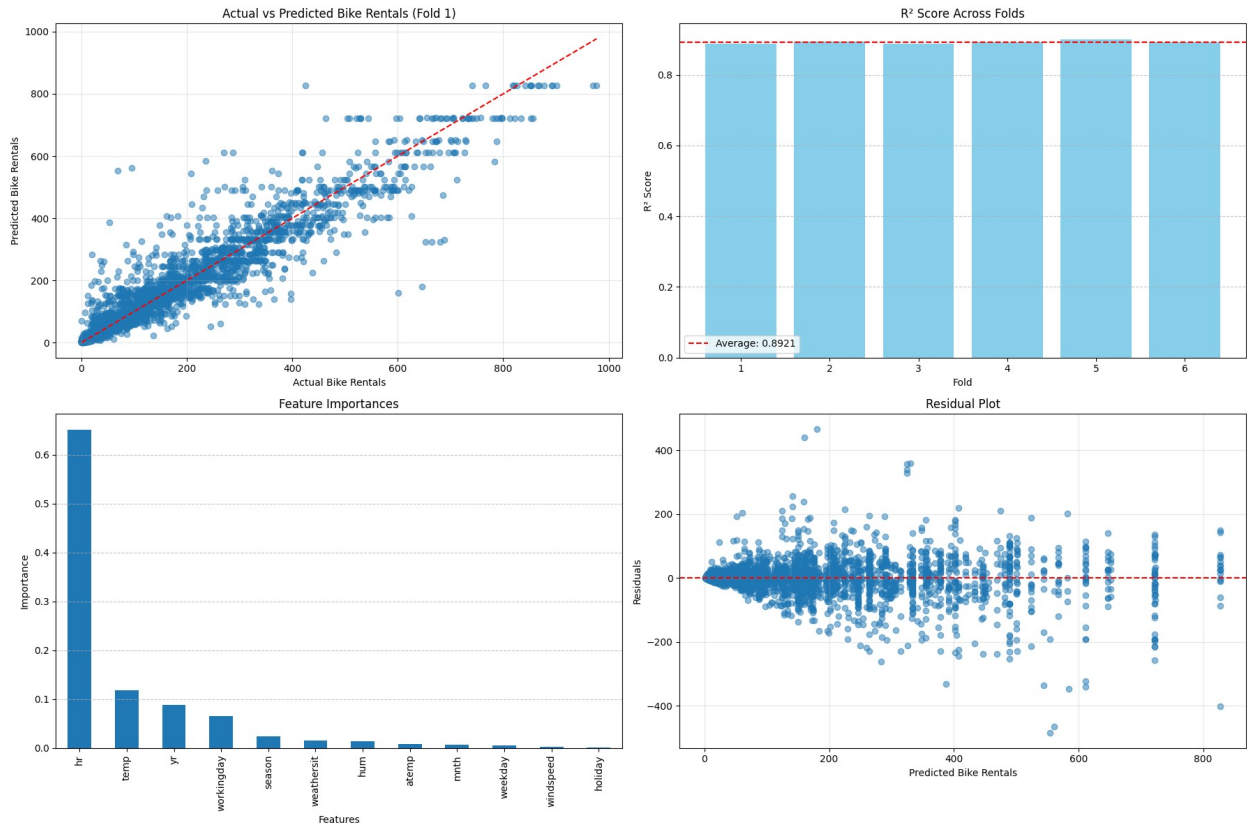
```

## Part 6: Decision Tree Regressor

SINGLE FOLD RESULTS (Fold 1):  
R<sup>2</sup> Score: 0.8872

Mean Squared Error: 3560.2039  
Mean Absolute Error: 35.5797  
Runtime: 0.0330 seconds

OVERALL CROSS-VALIDATION RESULTS:  
Average  $R^2$  Score: 0.8921 ( $\pm 0.0041$ )  
Average MSE: 3543.3972 ( $\pm 98.4431$ )  
Average MAE: 36.1249 ( $\pm 0.3392$ )  
Average Runtime: 0.0283 seconds



#### DECISION RULES EXTRACTED (First Fold):

Rule 1: IF (hr  $\leq$  -0.731) AND (hr  $\leq$  -0.876) AND (hr  $\leq$  -1.454) AND (workingday  $\leq$  -0.386) AND (atemp  $\leq$  0.008) AND (season  $\leq$  -0.008) AND (yr  $\leq$  -0.005) AND (atemp  $\leq$  -0.784) AND (holiday  $\leq$  2.814) AND (atemp  $\leq$  -2.366) THEN value = 12.33 [samples = 3]  
Rule 2: IF (hr  $\leq$  -0.731) AND (hr  $\leq$  -0.876) AND (hr  $\leq$  -1.454) AND (workingday  $\leq$  -0.386) AND (atemp  $\leq$  0.008) AND (season  $\leq$  -0.008) AND (yr  $\leq$  -0.005) AND (atemp  $\leq$  -0.784) AND (holiday  $\leq$  2.814) AND (atemp  $>$  -2.366) THEN value = 26.76 [samples = 45]  
Rule 3: IF (hr  $\leq$  -0.731) AND (hr  $\leq$  -0.876) AND (hr  $\leq$  -1.454) AND (workingday  $\leq$  -0.386) AND (atemp  $\leq$  0.008) AND (season  $\leq$  -0.008) AND (yr  $\leq$  -0.005) AND (atemp  $\leq$  -0.784) AND (holiday  $>$  2.814) AND (temp

```

<= -1.488) THEN value = 16.50 [samples = 2]
Rule 4: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr <= -0.005) AND (atemp <= -0.784) AND (holiday > 2.814) AND (temp >
-1.488) THEN value = 11.33 [samples = 3]
Rule 5: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr <= -0.005) AND (atemp > -0.784) AND (season <= -0.911) AND (hum <=
1.027) THEN value = 37.62 [samples = 8]
Rule 6: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr <= -0.005) AND (atemp > -0.784) AND (season <= -0.911) AND (hum >
1.027) THEN value = 17.00 [samples = 2]
Rule 7: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr <= -0.005) AND (atemp > -0.784) AND (season > -0.911) AND (hum <=
-0.244) THEN value = 90.00 [samples = 2]
Rule 8: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr <= -0.005) AND (atemp > -0.784) AND (season > -0.911) AND (hum > -
0.244) THEN value = 46.56 [samples = 9]
Rule 9: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr > -0.005) AND (temp <= -0.659) AND (hum <= -1.048) AND (windspeed
<= 0.399) THEN value = 91.00 [samples = 2]
Rule 10: IF (hr <= -0.731) AND (hr <= -0.876) AND (hr <= -1.454) AND
(workingday <= -0.386) AND (atemp <= 0.008) AND (season <= -0.008) AND
(yr > -0.005) AND (temp <= -0.659) AND (hum <= -1.048) AND (windspeed
> 0.399) THEN value = 66.14 [samples = 7]
... [total of 899 rules]

DecisionTreeRegressor(max_depth=3, random_state=42)

```

## Comments-6

### 1. Problem and Approach

In this part, I built a **Decision Tree Regressor** using the **Bike Sharing dataset**.

Decision trees are useful for regression because they can **model non-linear relationships** and **capture feature interactions** effectively.

- I used **K-Fold Cross-Validation (6 folds)** to ensure a fair evaluation.
- I have applied **two different pruning strategies**:
  - **Pre-Pruning (max\_depth, min\_samples\_split)** to prevent overfitting.
  - **Post-Pruning (Cost Complexity Pruning - CCP)** to remove unnecessary branches and improve generalization.

---

## 2. Dataset and Preprocessing

- The dataset contains daily **bike rental counts** with various weather and time-related features.
  - The target variable (`cnt`) represents the **total number of rentals per day**.
  - **Feature scaling** was applied to ensure consistent distance-based calculations.
- 

## 3. K-Fold Cross Validation

To improve the reliability of the evaluation, I used **6-fold cross-validation**:

- The dataset was split into **6 equal parts**, with each part used as a test set once.
  - This method helped to **reduce overfitting** and provided a **more stable performance estimate**.
- 

## 4. Performance Metrics

Since this is a regression task, I evaluated the model using:

- **R<sup>2</sup> Score** → Measures how well the model explains the data (**closer to 1 is better**).
  - **Mean Squared Error (MSE)** → Measures how large the errors are (**smaller is better**).
  - **Mean Absolute Error (MAE)** → Measures the average absolute error (**smaller is better**).
  - **Runtime** → Checked to see how efficient the model is.
- 

## 5. Results and Interpretation

- The model achieved an **average R<sup>2</sup> score of 0.8921**, meaning it explains **about 89% of the variance** in bike rentals.
  - **MSE (3543.40) and MAE (36.12)** indicate that the model makes reasonable predictions, but there is still some error.
  - **The runtime was very low (~0.03 seconds)**, meaning the model is computationally efficient.
  - The **performance across folds was consistent** (small standard deviation), indicating a **stable model**.
-

## 6. Comparison with Other Regressors

- **Decision Tree vs. KNN (Part 2):**
  - Decision Tree **achieved the highest  $R^2$  score (89.21%)**, meaning it captured the most variance in bike rentals.
  - KNN **performed well (84.12%)** but is **slower**, as it needs to compute distances for all data points.
  - **Decision Tree is interpretable**, while KNN acts as a black-box model.
- **Decision Tree vs. SVM (Part 4):**
  - Decision Tree **outperformed SVM (81.45%)**, showing it handles **complex relationships better**.
  - SVM assumes a **linear relationship**, which might not be ideal for this dataset.
  - Decision Trees **are faster** and **do not require feature scaling**, unlike SVM.